

Section 1 - Computer Arithmetic & Numerical Accuracy

M1 Mandelbrot Trajectory Divergence

To be able to see the impact of the numerical accuracy a trajectory divergence is performed on the Mandelbrot set using the complex128 and complex64 data types. The trajectory divergence test was done in the Seahorse Valley, a region where the complex constant C is initialized as $R = [-0.753; -0.749]$ and $I = [0.099; 0.103]$. The test was performed with a Mandelbrot set with the resolution being 1024x1024, with the maximum number of iterations being 1000, and divergence threshold $\tau = 0.01$. The code can be found in `mandelbrot_trajectory_divergence.py`.

What is the fraction of pixels that diverge before max iter?

I find that the fraction of pixels diverging before the maximum number of iterations is 0.9999981, so nearly all pixels diverge before reaching the maximum number of iterations. In total only two pixels reaches the maximum number of iterations.

If we instead set the maximum number of iterations to 100, the fraction becomes 0.9846.

Where do trajectories diverge early? Compare visually to the escape-count map.

By looking at just Figure 1, it seems that the trajectory diverges early in exterior areas of the Mandelbrot set, but also in some of the interior areas. However, looking at Figure 2, it can be clearly seen that this is not the case, as the maximum iteration before divergence is consistently reached across the interior of the set. This means that in the exterior of the set and the border areas the numerical accuracy quickly deteriorates.




Figure 1, (Left) The Seahorse Valley of the Mandelbrot set. (Right) The trajectory divergence map of the Seahorse Valley.




Figure 2, Divergence map for the entire Mandelbrot Set.

Does early divergence correlate with high escape iteration counts?

No, it does not. It is mostly the opposite, which means high escape iteration count correlate with late divergence and vice versa.

M2 Mandelbrot Sensitivity map

The sensitivity map is computed by calculating κ , which approximates the relative change in input to the relative change in output.

$$\kappa(c) = \frac{|\Delta n|}{\epsilon_{32} \cdot n(c)}$$

$$\Delta n = n(c + \delta) - n(c)$$

$$\delta = \epsilon_{32} \cdot |c|$$

c is a complex constant.

ϵ_{32} is floating-point accuracy, defined as the difference between 1.0 and the next float of the same size that can be represented.

δ is a perturbation to the complex constant.

$n(c)$ is the number of iterations it takes before a pixel escapes.

Δn is the change in the number of iterations before a pixel escape, when adding the perturbation δ to the complex constant c and subtracting that from $n(c)$.




Figure 3, Sensitivity map for: (Left) Seahorse Valley, (Right) The entire Mandelbrot set.

Where is κ largest? Does it match the boundary in M1?

When looking at Figure 3 left, κ seems to be largest near the boundaries and in the interior. However, when looking at Figure 3 right, it can be clearly seen that κ is largest in the boundary areas and stable elsewhere. This matches with M1, except for the exterior. This may be because the complex number in the exterior quickly grows beyond the value it would have escaped with in the normal Mandelbrot Set.

What is κ for interior pixels ($n = \max \text{ iter}$)?

Again, when looking at Figure 3 left, it is difficult to determine. However, on Figure 3 right, κ is clearly small for the interior pixels. This means that the problem is well-conditioned only in the non-boundary. Whereas the problem is ill-conditioned in the boundary areas, which means that small input perturbations δ propagate and become large output errors κ .

Section 2 - Testing & Documentation

M1 Test Suite

I created a test suite consisting of 6 unit tests. All of the tests passed after I fixed a few undiscovered bugs that were visually undetectable. The coverage of the unit tests was quite low at only 11%. The reason for this low coverage is because my project folder contains both test files as well as exercise files, such as the Monte Carlo pi estimation, as seen on Figure 4. These files do not need testing, as they are not part of the Mandelbrot project, making the coverage quite meaningless in this case.

```

===== test session starts =====
platform win32 -- Python 3.11.14, pytest-9.0.2, pluggy-1.6.0 -- C:\Users\srisb\local\share\mamba\envs\nsc2026\python.exe
cachedir: .pytest_cache
rootdir: C:\Users\srisb\OneDrive\Dokumenter\Alt(Rog)\Universitet\8.Semester\Numerical Scientific Computing\nsc-course
plugins: cov-7.1.0
collected 8 items

test_mandelbrot.py::test_numba_mandelbrot_point[0j-100] PASSED [ 12%]
test_mandelbrot.py::test_numba_mandelbrot_point[(-1+1j)-2] PASSED [ 25%]
test_mandelbrot.py::test_numba_mandelbrot_point[(-1.5+0.01j)-13] PASSED [ 37%]
test_mandelbrot.py::test_mandelbrot_mp_chunk_output PASSED [ 50%]
test_mandelbrot.py::test_mandelbrot_mp_chunk_output_length PASSED [ 62%]
test_mandelbrot.py::test_mandelbrot_vec_correctness PASSED [ 75%]
test_mandelbrot.py::test_mandelbrot_numba_correctness PASSED [ 87%]
test_mandelbrot.py::test_numerical_differences PASSED [100%]

===== tests coverage =====
coverage: platform win32, python 3.11.14-final-0

Name                               Stmts  Miss  Cover
-----
dask_test.py                        7      5    29%
julia.py                           39     39     0%
line_profile.py                    16     16     0%
mandelbrot_dask.py                  74     74     0%
mandelbrot_dask_distributed.py     104    104     0%
mandelbrot_dask_distributed_sweep.py 186    186     0%
mandelbrot_dask_sweep.py           184    184     0%
mandelbrot_mp.py                    90     77    14%
mandelbrot_mp_chunked.py           97     97     0%
mandelbrot_naive.py                 47     28    57%
mandelbrot_numba.py                 50     48    20%
mandelbrot_numba_parallel.py        49     49     0%
mandelbrot_sensitivity_map.py       36     36     0%
mandelbrot_time_scaling.py          18     18     0%
mandelbrot_trajectory_divergence.py 37     37     0%
mandelbrot_vectorized.py            41     21    49%
memory_layout_performance.py        16     16     0%
memory_type_performance.py          21     21     0%
monte_carlo_pi_dask.py              33     33     0%
monte_carlo_pi_est.py               52     52     0%
numba_benchmark.py                  22     22     0%
profiling.py                        21     21     0%
test_install.py                     13      0   100%
test_mandelbrot.py                  46      1    98%
TOTAL                             1139   1009    11%

8 passed in 120.68s (0:02:00)

```

Figure 4 pytest test coverage.

M2 Docstrings and Type Hints

Docstrings on the NumPy style format was added to `mandelbrot_naive.py` along with type hints for both parameters and return variables. A slight problem with type hinting NumPy arrays is that NumPy does not specify how to type hint the shape an array. It is possible to type hint the array type, but as this can also change, I have just used the basic `np.ndarray` type. All ruff checks were passed as shown in Figure 5.

```

(nsc2026) C:\Users\srisb\OneDrive\Dokumenter\Alt(Rog)\Universitet\8.Semester\Numerical Scientific Computing\nsc-course>ruff check mandelbrot_naive.py
All checks passed!

```

Figure 5 Verification that all ruff checks were passed in `mandelbrot_naive.py`.

Section 3 – GPU Computing

M1 Mandelbrot float 32 kernel

My Mandelbrot kernel uses the 32-bit floating-point data type, and can compute the Mandelbrot set with a resolution of 8192x8192 with 100 maximum iterations in 0.0182 seconds on my Nvidia RTX 2060 max-q. This timing is only for the execution of the kernel and is found using the median of 50 runs to ensure consistency. If I account for copying data to and from the kernel, the Mandelbrot set is computed in 0.1119 seconds, which shows that the memory transfers from host to device and vice versa has a huge impact on the runtime.

M2 Mandelbrot float32 vs float64

As seen in M1 when executing the kernel using 32-bit floating points it takes 0.1119 seconds to compute the Mandelbrot set with a resolution of 8192x8192 and 100 maximum iterations. When using 64-bit floating points instead the computation takes 0.6745 seconds instead.

Is the measured speed ratio consistent with your Roofline prediction?

By looking at [Tech Power Up](#) website the specs for my Nvidia RTX 2060 max-q can be found. We can now calculate the ridge point for both float32 and float64.

$$bandwidth = 264.0 \text{ GB/s}$$

$$perf_{32} = 4.550 \text{ TFLOPS}$$

$$perf_{64} = 142.2 \text{ GFLOPS}$$

Where FLOPS are floating point operations per second. We can now calculate the ridge point as:

$$ridge = \frac{perf}{bandwidth}$$

We calculate and get:

$$ridge_{32} = \frac{perf_{32}}{bandwidth} = \frac{4550.0 \text{ GFLOPS}}{264.0 \text{ GB/s}} = 16.856 \frac{\text{FLOPs}}{\text{byte}}$$

$$ridge_{64} = \frac{perf_{64}}{bandwidth} = \frac{142.2 \text{ GFLOPS}}{264.0 \text{ GB/s}} = 0.539 \frac{\text{FLOPs}}{\text{byte}}$$

We now see that theoretically the Nvidia RTX 2060 max-q should be ~31.27 slower when using the float64 datatype compared to the float32 one. We can now calculate the ratio and we get: $\frac{0.6745}{0.1119} = 6.0277$, which means we do not line up with the roofline prediction. The reason for this discrepancy is likely because we do not exclusively perform floating point operations in the kernel.

Another thing we can observe is that the problem is still compute bound; despite using a GPU, as the arithmetic intensity (45 FLOPs/byte (L3 p. 16)) is to the right of both ridge points. The runtimes can be found below in Table 1.

Implementation (8192x8192)	Time [s]
OpenCL numpy.float32	0.1119
OpenCL numpy.float64	0.6745

Table 1, runtimes for OpenCL using 32-bit and 64-bit floating point operations.

M3 Benchmarking

In Table 2 and Table 3 the different implementations' benchmark timings from previous lectures and assignments.

1024x1024	Time [s]	Speedup	Workers/cores
Naïve	39.858	1	1
Vectorized	1.2566	31.7	1
NJIT Hybrid	1.1809	33.8	1
NJIT (f32, f64)	0.1699	234.6	1
NJIT Parallel	0.0191	2086.8	8
Multiprocessing	0.0171	2330.9	16
Dask Local	0.1782	223.7	8
OpenCL float32	0.0045	8857.4	1920 ¹
OpenCL float64	0.0146	2730	1920

Table 2, Performance table for the Mandelbrot set with a resolution of 1024x1024.

8192x8192	Time [s]	Speedup	Workers/cores
NJIT Parallel	1.1353	1	1
Multiprocessing	0.5531	2.0526	16
Dask Local	2.7796	0.4084	8
Dask Distributed	1.7715	0.6409	16
OpenCL float32	0.1119	10.146	1920
OpenCL float64	0.6745	1.6832	1920

Table 3, Performance table for the Mandelbrot set with a resolution of 8192x8192.

The tables have also been visualized as bar charts with a logarithmic y-axis in Figure 6 and Figure 7.

¹ See the [Tech Power Up](#) website.




Figure 6, Bar plot of timings for 1024x1024 implementations.




Figure 7, Bar plot of timings for 8192x8192 implementations.

Where does GPU sit relative to Numba? To naive Python? (Recall the L01 starting point.)

The amount of compute that the GPU is capable of heavily outweighs the transfers from host to device and vice versa for a perfectly parallelizable problem, which computing the Mandelbrot set is. This is also clearly shown in Figure 6 and Figure 7 where the GPU implementations beat all other implementations when using the 32-bit floating point data type.

Reflections MP3

Reflections have been added here for completeness.

When is the GPU the right tool — and when is it not worth the effort compared to Numba or multiprocessing?

The GPU is the right tool when the problem is very parallelizable, such as computing the Mandelbrot set. However, when there is overhead in transferring data from the host to the device and vice versa. Therefore, if the problem is not large enough, it might not be worth it. Another problem with using the GPU is the possible complexity of implementing an OpenCL or CUDA kernel, which almost always will take longer to implement than simply decorating a python function with @NJIT.

If the work that must be performed by each core contains many branch statements, it is better to use multiprocessing or NJIT, than the GPU.

Did writing tests reveal anything unexpected about your code? Did you need to refactor anything to make it testable?

Yes, as mentioned in my MP3.pdf I discovered that the numerical accuracy of the Mandelbrot set was not equal across implementations, which shows why you should write the tests first and then implement the code that should pass the tests.

After adding docstrings and type hints: does the effort feel worthwhile? What would have been easier if you had written them from the start?

I like programming in C/C++ much more because types are required. Therefore, when programming in Python I always add type hints because it makes the code so much more readable and therefore worthwhile. However, there are problems with type hints especially when using NumPy as there are no standard for defining array shapes. Docstrings are useful especially in large projects, but in this course I only benefitted very little from them.